

lamprey_rnaseq_sample_comp

April 30, 2014

1 Lamprey RNA-Seq 2014

1.1 Between-sample Assembly Comparison

1.1.1 Camille Scott (camille.scott.w@gmail.com) Michigan State University GED Lab

1.1.2 <https://github.com/camillescott/2013-lamprey/tree/lamp3>

1.2 Overview

The primary goal of this project has been to produce a scientifically usable set of transcripts from our vast collection of RNA-seq samples. Given the lack of completeness of the current lamprey genome, we chose to assemble the reads de novo, using our in-house pipeline. The basics of the assembly pipeline can be found in the [Eel Pond protocol](#).

This document goes over the content of the resulting assembly with respect to each sample used to construct it. We align every sample to the assembly, and use an abundance estimation tool to get expression levels for each transcript, for each sample. This allows us to analyze the extent to which each sample contributed to the final assembly, and the relative representation of each sample within the assembly. We consider a transcript *represented* in a sample if it has an estimated abundance greater than 1.0 from the reads in that sample.

An overview of the samples can be found [here](#). This is what I have been able to scrape up from various collaborators and sequencing centers; some information is definitely missing, and some samples were preprocessed before being given to me. As such, any comments on this document are greatly appreciated, whether they be corrections or offers for original versions of the data. Note that the second sheet has the key between the original sample names and the id's being used in the current project.

For all plots, “filtered” means after the removal of known contaminants.

All the plots shown in this notebook, as well as many more supporting figures, can be found in the full notebook in the previously linked github repository, and can be viewed in rendered form [here](#).

1.3 Validation

Briefly, we can show that the assembly is both valid and contains new information with homologies to model organisms. Here, we show:

1. **Putative new:** Transcripts with substantial protein homology to a model species (mouse, zebrafish, or amphioxus) or substantial nucleotide homology to the lamprey genome, but no homology with the existing lamprey transcripts or ESTs.
2. **All putative:** Transcripts which have homologies in any one of the model species, the lamprey genome, lamprey transcripts, or lamprey ESTs.

This is done for the transcripts represented in each sample:

Figure 1: filtered_putative_transcripts

Note that there is overlap between bars; that is, the sum under the “curve” is much greater than the total number of transcripts in the assembly.

With that in mind, the distribution of the number of samples in which a transcript is represented, along with genome homologies for each bin:

Figure 2: transcript_sample_dist_support

1.4 Sample Comparison

Between sample comparisons seek to analyze both the transcript content and transcript expression of different samples. The manner in which these metrics is analyzed depends mostly on the *distance* metric used.

Briefly, let us have an assembly T , consisting of transcripts $T = \{t_0, t_1, \dots, t_n\}$, where there are n transcripts total. Let us also have a set of samples S , where $S = \{s_0, s_1, \dots, s_m\}$, with m samples total; In our case, $m = 84$.

Then each sample s_i is a vector $\bar{s}_i = [A_i(t_0), A_i(t_1), \dots, A_i(t_n)]^T$, where

$$A_i(t_j) = \begin{cases} 1, & \text{if estimated abundance of transcript } t_j \text{ in sample } i \text{ is } > 1.0 \\ 0, & \text{otherwise} \end{cases}$$

for simple sample transcript content, and

$$A_i(t_j) = \text{estimated abundance of transcript } t_j \text{ in sample } i$$

for sample expression.

We then need only define a distance metric $D(s_a, s_b)$ to compare two samples. This facilitates clustering and tree-building between all pairs of samples.

In qualitative terms, we simply have 84 lists, 1 for each sample, each of which is a list of estimated abundances for each transcript. We’ll do pairwise comparisons of these lists.

The other important choice is the clustering method. Two methods will be demonstrated: the *centroid* method and the *ward* method. There are decent descriptions of these in the [scipy clustering docs](#), with more detailed descriptions on wikipedia and in many textbooks and papers. Briefly, in the *centroid* method, the distance between two clusters is the distance between the mean positions of all elements in the clusters. It is noted as being robust to outliers, but potentially performing worse than other methods. In the *Ward* method, the distance between two clusters is the *ANOVA* sum of squares between the two clusters; we seek to minimize the within-cluster sum of squares distance. *Ward’s* method performs well with small N , and is strongly biased toward producing clusters with the same number of observations. See [here](#) for these overviews.

1.4.1 Sample Content: *Hamming Distance*

Hamming distance measures the number of disagreeing positions in a binary vector, normalized by the length of the vector (number of transcripts in this case). See the [scipy docs](#) for a detailed description of the metric used in these figures.

Qualitatively, this is a measure of how many transcripts are shared between the two samples.

Figure 3: sample_dendro_hamming_centroid_filtered

Figure 4: sample_dendro_hamming_ward_filtered

Here, the *Ward* method appears to give better results.

1.4.2 Sample Expression: *Correlation Distance*

Correlation distance is a measure of the statistical independence of two distributions. Thus, it will take into account the estimated abundances, rather than just considering if the transcripts are represented at all, as in the Hamming distance. The [scipy docs](#) once more have a good description of what was used for this figure; further information is available on [wikipedia](#). Notably, while classically a correlation distance of 0.0 implies complete statistical independence, this implementation reports $1 - CD(s_a, s_b)$; ie, 1.0 implies complete statistical independence.

Qualitatively, this is a measure of how similar the distribution of expression levels (estimated abundances) is between two samples.

Figure 5: sample_dendro_correlation_centroid_filtered

Figure 6: sample_dendro_correlation_ward_filtered

1.4.3 Sample Expression: *Specific Cases*

We can also look at how the clustering changes with differing subsets of transcripts. This could be of interest in regards to evolutionary changes.

First, we will take transcripts homologous to an amphioxus protein, but not to mouse or zebrafish. Note that amphioxus shares a common ancestor with lamprey.

Figure 7: amphioxus_dendro_hamming_centroid

Next, we take transcripts homologous to mouse or zebrafish, but not to amphioxus. Note that mouse and zebrafish and descendants of lamprey.

1.5 Conclusions

For now, I'm leaving out any of my interpretations of the data. I welcome any and all feedback, thoughts, ideas – especially any thoughts on how this analysis (and other similar analyses) that I can perform are useful for your research.

Figure 8: amphioxus_dendro_correlation_centroid

Figure 9: zebrafish_dendro_hamming_centroid

Figure 10: zebrafish_dendro_correlation_centroid